

## lab2\_finale

June 4, 2025

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: train = pd.read_csv('mai-ml-lab-2/train.csv')
test = pd.read_csv('mai-ml-lab-2/test.csv')

display(train.info())
display(test.info())

print('Train duplication: ', train.duplicated().sum())
print('Test duplication: ', test.duplicated().sum())
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 15000 entries, 0 to 14999

Data columns (total 24 columns):

#	Column	Non-Null Count	Dtype
0	id	15000 non-null	int64
1	age	15000 non-null	float64
2	height(cm)	15000 non-null	float64
3	weight(kg)	15000 non-null	float64
4	waist(cm)	15000 non-null	float64
5	eyesight(left)	15000 non-null	float64
6	eyesight(right)	15000 non-null	float64
7	hearing(left)	15000 non-null	float64
8	hearing(right)	15000 non-null	float64
9	systolic	15000 non-null	float64
10	relaxation	15000 non-null	float64
11	fasting blood sugar	15000 non-null	float64
12	Cholesterol	15000 non-null	float64
13	triglyceride	15000 non-null	float64
14	HDL	15000 non-null	float64
15	LDL	15000 non-null	float64
16	hemoglobin	15000 non-null	float64
17	Urine protein	15000 non-null	float64
18	serum creatinine	15000 non-null	float64

```

19  AST                15000 non-null float64
20  ALT                15000 non-null float64
21  Gtp                15000 non-null float64
22  dental caries      15000 non-null float64
23  smoking            15000 non-null float64

```

dtypes: float64(23), int64(1)

memory usage: 2.7 MB

None

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 10000 entries, 0 to 9999

Data columns (total 23 columns):

#	Column	Non-Null Count	Dtype
0	id	10000 non-null	int64
1	age	10000 non-null	float64
2	height(cm)	10000 non-null	float64
3	weight(kg)	10000 non-null	float64
4	waist(cm)	10000 non-null	float64
5	eyesight(left)	10000 non-null	float64
6	eyesight(right)	10000 non-null	float64
7	hearing(left)	10000 non-null	float64
8	hearing(right)	10000 non-null	float64
9	systolic	10000 non-null	float64
10	relaxation	10000 non-null	float64
11	fasting blood sugar	10000 non-null	float64
12	Cholesterol	10000 non-null	float64
13	triglyceride	10000 non-null	float64
14	HDL	10000 non-null	float64
15	LDL	10000 non-null	float64
16	hemoglobin	10000 non-null	float64
17	Urine protein	10000 non-null	float64
18	serum creatinine	10000 non-null	float64
19	AST	10000 non-null	float64
20	ALT	10000 non-null	float64
21	Gtp	10000 non-null	float64
22	dental caries	10000 non-null	float64

dtypes: float64(22), int64(1)

memory usage: 1.8 MB

None

Train duplication: 0

Test duplication: 0

```

[3]: print(f"Train shape: {train.shape}")
      print(f"Test shape: {test.shape}")
      print("Columns in train:", train.columns.tolist())

```

```

Train shape: (15000, 24)
Test shape: (10000, 23)
Columns in train: ['id', 'age', 'height(cm)', 'weight(kg)', 'waist(cm)',
'eyesight(left)', 'eyesight(right)', 'hearing(left)', 'hearing(right)',
'systolic', 'relaxation', 'fasting blood sugar', 'Cholesterol', 'triglyceride',
'HDL', 'LDL', 'hemoglobin', 'Urine protein', 'serum creatinine', 'AST', 'ALT',
'Gtp', 'dental caries', 'smoking']

```

```

[4]: features = list(set(train.columns) - {'id', 'smoking'})

df_all = pd.concat([train[features], test[features]], axis=0)

```

```

[5]: print('Duplicated rows between train and test datasets: ', df_all.duplicated().
      ↪sum())

```

Duplicated rows between train and test datasets: 0

```

[6]: display(train.head(5))
train.info()
display(train.describe())

```

	id	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	\
0	0	55.0	150.0	50.0	73.0	0.8	
1	1	50.0	160.0	60.0	79.6	1.0	
2	2	55.0	145.0	50.0	73.0	1.2	
3	3	40.0	175.0	80.0	83.2	0.9	
4	4	30.0	170.0	70.0	79.5	1.2	

	eyesight(right)	hearing(left)	hearing(right)	systolic	...	HDL	LDL	\
0	1.0	1.0	1.0	122.0	...	62.0	113.0	
1	0.8	1.0	1.0	124.0	...	74.0	112.0	
2	1.2	1.0	1.0	110.0	...	66.0	123.0	
3	1.0	1.0	1.0	99.0	...	38.0	114.0	
4	1.2	1.0	1.0	119.0	...	49.0	148.0	

	hemoglobin	Urine protein	serum creatinine	AST	ALT	Gtp	\
0	12.0	1.0	0.6	22.0	21.0	12.0	
1	14.1	1.0	0.9	19.0	14.0	17.0	
2	13.3	1.0	0.6	15.0	12.0	13.0	
3	16.1	1.0	1.0	22.0	13.0	20.0	
4	15.5	1.0	1.2	18.0	19.0	25.0	

	dental caries	smoking
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	1.0
4	0.0	1.0

[5 rows x 24 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 15000 entries, 0 to 14999

Data columns (total 24 columns):

#	Column	Non-Null Count	Dtype
0	id	15000 non-null	int64
1	age	15000 non-null	float64
2	height(cm)	15000 non-null	float64
3	weight(kg)	15000 non-null	float64
4	waist(cm)	15000 non-null	float64
5	eyesight(left)	15000 non-null	float64
6	eyesight(right)	15000 non-null	float64
7	hearing(left)	15000 non-null	float64
8	hearing(right)	15000 non-null	float64
9	systolic	15000 non-null	float64
10	relaxation	15000 non-null	float64
11	fasting blood sugar	15000 non-null	float64
12	Cholesterol	15000 non-null	float64
13	triglyceride	15000 non-null	float64
14	HDL	15000 non-null	float64
15	LDL	15000 non-null	float64
16	hemoglobin	15000 non-null	float64
17	Urine protein	15000 non-null	float64
18	serum creatinine	15000 non-null	float64
19	AST	15000 non-null	float64
20	ALT	15000 non-null	float64
21	Gtp	15000 non-null	float64
22	dental caries	15000 non-null	float64
23	smoking	15000 non-null	float64

dtypes: float64(23), int64(1)

memory usage: 2.7 MB

	id	age	height(cm)	weight(kg)	waist(cm)	\
count	15000.000000	15000.000000	15000.000000	15000.000000	15000.000000	
mean	7499.500000	42.606000	164.672667	64.282000	80.323180	
std	4330.271354	10.415341	8.558684	10.973871	7.932701	
min	0.000000	20.000000	140.000000	35.000000	0.900000	
25%	3749.750000	40.000000	160.000000	55.000000	75.000000	
50%	7499.500000	40.000000	165.000000	65.000000	80.000000	
75%	11249.250000	50.000000	170.000000	70.000000	86.000000	
max	14999.000000	80.000000	190.000000	120.000000	121.000000	

	eyesight(left)	eyesight(right)	hearing(left)	hearing(right)	\
count	15000.000000	15000.000000	15000.000000	15000.000000	
mean	1.032727	1.025927	1.005533	1.004933	
std	0.312845	0.313488	0.074183	0.070067	
min	0.100000	0.100000	1.000000	1.000000	

25%	0.800000	0.800000	1.000000	1.000000
50%	1.000000	1.000000	1.000000	1.000000
75%	1.200000	1.200000	1.000000	1.000000
max	9.900000	9.900000	2.000000	2.000000

	systolic	...	HDL	LDL	hemoglobin	\
count	15000.000000	...	15000.000000	15000.000000	15000.000000	
mean	119.149467	...	57.535667	113.158533	14.510340	
std	11.167046	...	11.380739	18.442012	1.392921	
min	82.000000	...	28.000000	41.000000	7.200000	
25%	110.000000	...	49.000000	100.000000	13.500000	
50%	119.000000	...	56.000000	112.000000	14.700000	
75%	128.000000	...	65.000000	126.000000	15.600000	
max	190.000000	...	105.000000	193.000000	28.900000	

	Urine protein	serum creatinine	AST	ALT	\
count	15000.000000	15000.000000	15000.000000	15000.000000	
mean	1.016533	0.866580	22.272667	20.874933	
std	0.157464	0.168492	6.324422	10.314009	
min	1.000000	0.100000	9.000000	4.000000	
25%	1.000000	0.800000	18.000000	14.000000	
50%	1.000000	0.900000	21.000000	18.000000	
75%	1.000000	1.000000	25.000000	24.000000	
max	5.000000	1.600000	324.000000	190.000000	

	Gtp	dental caries	smoking
count	15000.000000	15000.000000	15000.000000
mean	26.290200	0.139467	0.366600
std	20.431882	0.346444	0.481892
min	6.000000	0.000000	0.000000
25%	15.000000	0.000000	0.000000
50%	21.000000	0.000000	0.000000
75%	31.000000	0.000000	1.000000
max	566.000000	1.000000	1.000000

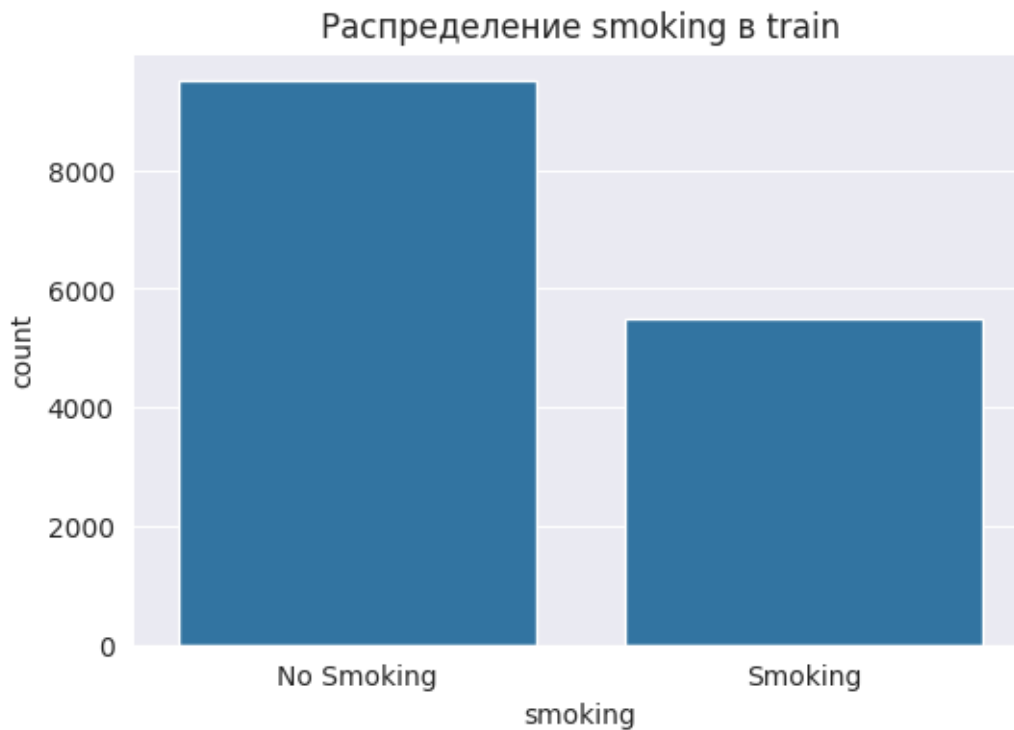
[8 rows x 24 columns]

```
[7]: missing = train.isnull().sum().sort_values(ascending=False)
print("      train:")
print(missing[missing > 0])
```




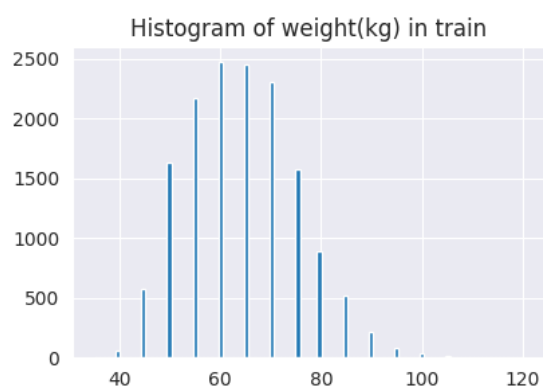


```
      train:
Series([], dtype: int64)
```







```
[8]: plt.figure(figsize=(6,4))
sns.countplot(x='smoking', data=train)
plt.title('      smoking      train')
plt.xticks(ticks=[0, 1], labels=['No Smoking', 'Smoking'])
```

```
plt.ylabel('count')
plt.show()
```





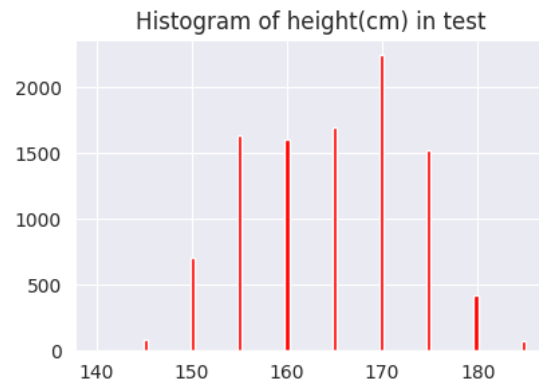







```
[9]: def features_histogram(train, test):
    for feature in features:
        plt.figure(figsize=(10, 3))
        plt.subplot(1, 2, 1)
        train[feature].hist(bins=100)
        plt.title(f"Histogram of {feature} in train")
        plt.subplot(1, 2, 2)
        test[feature].hist(bins=100, color='r')
        plt.title(f"Histogram of {feature} in test")
        plt.show()
    features_histogram(train, test)
```




















































```
[10]: for feature in features:
        feature_min = test[feature].min()
        feature_max = test[feature].max()
        # adjust train data
        train = train[(train[feature] >= feature_min) & (train[feature] <=
        feature_max)]







train = train.reset_index(drop=True)
features_histogram(train, test)
```










































### 0.0.1 Checking statistic of train and test

```
[11]: plt.figure(figsize=(7, 7))

plt.plot(train[features].mean(), test[features].mean(), 'r.', label='Mean')
plt.plot(train[features].std(), test[features].std(), 'g.', label='Deviation')
plt.plot(train[features].skew(), test[features].skew(), 'b.', label='Skewness')
plt.plot(train[features].kurtosis(), test[features].kurtosis(), 'k.', label='Kurtosis')


plt.plot(train[features].min(), test[features].min(), 'c.', label='Min')
plt.plot(train[features].max(), test[features].max(), 'm.', label='Max')

plt.plot([1e-2, 1e4], [1e-2, 1e4], 'y')
plt.title('Comparison of Train and Test Statistics')
plt.xlabel('Train Statistics')
plt.ylabel('Test Statistics')

plt.xlim(1e-2, 1e4)
plt.ylim(1e-2, 1e4)
plt.xscale('log')
plt.yscale('log')

plt.legend()

plt.show()
```



## 0.0.2


```
[12]: train_corr = train[features].corr()
      test_corr = test[features].corr()

      plt.figure(figsize=(15, 15))


      sns.heatmap(train_corr, annot=True, cbar=False)
      plt.title('Correlation Matrix for Train')
      plt.show()

      plt.figure(figsize=(15, 15))
```

```
sns.heatmap(test_corr, annot=True, cbar=False)
plt.title('Correlation Matrix for Test')
plt.show()
```








### 0.0.3 Linear Correlation between test and train features

```
[13]: plt.plot(np.array(train_corr).flatten(), np.array(test_corr).flatten(), 'k.')
plt.plot([-0.5, 1], [-0.5, 1], 'y')
plt.xlim(-0.5, 1)
plt.ylim(-0.5, 1)

plt.xlabel('Train feature Correlation')
plt.ylabel('Test feature Correlation')

plt.title('Correlation comparison between Train and Test')
plt.show()
```



```
[14]: sorted(train[features].nunique())
```

```
[14]: [2,  
      2,  
      2,  
      4,  
      10,  
      13,  
      13,  
      14,  
      14,  
      15,  
      57,  
      62,  
      74,  
      75,  
      96,  
      107,  
      109,  
      133,
```

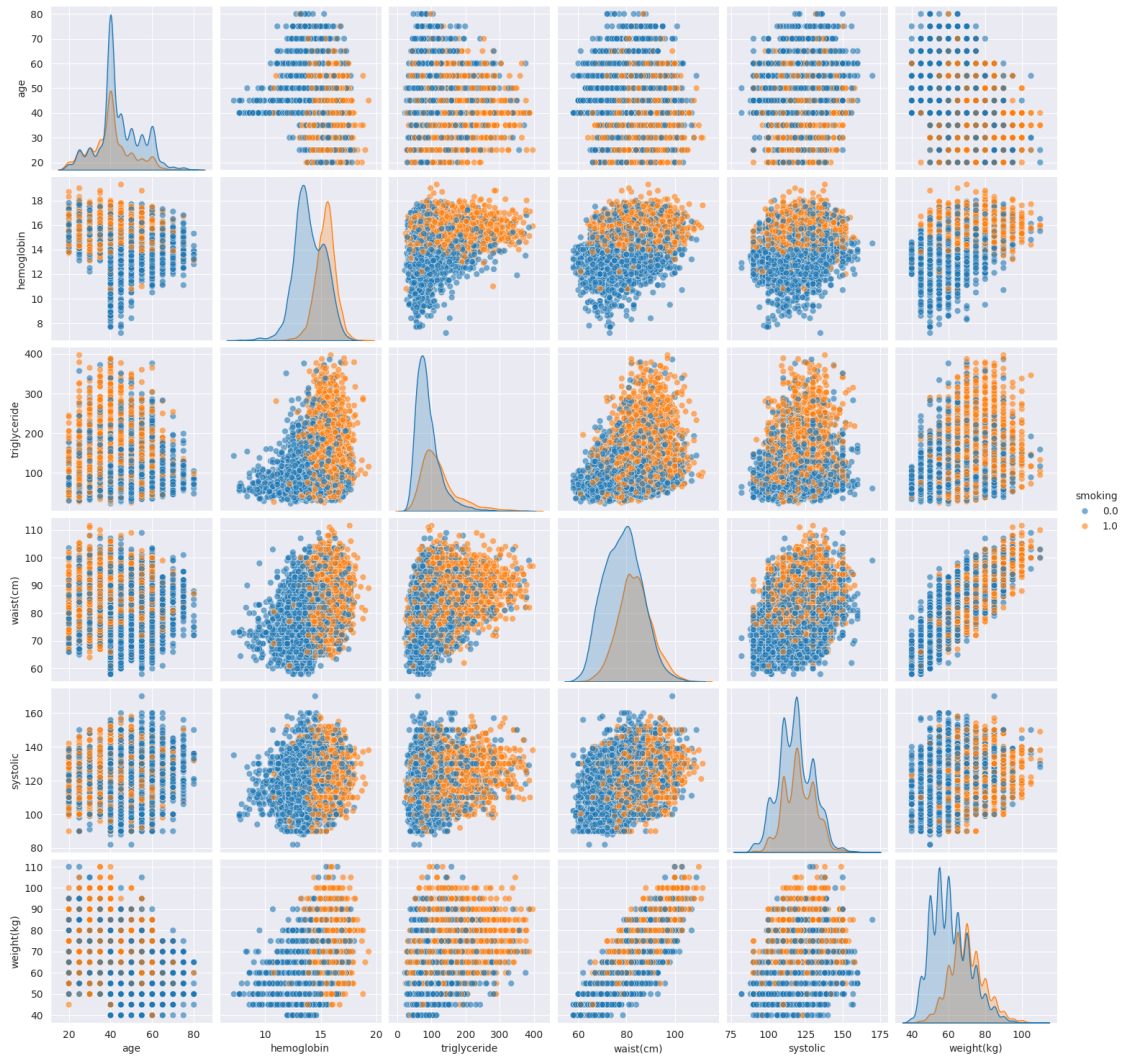
```
148,  
168,  
327,  
353]
```


#### 0.0.4


```
[15]: overlooked_features = ['age', 'hemoglobin', 'triglyceride', 'waist(cm)',  
    ↪ 'systolic', 'weight(kg)']  
  
sns.pairplot(  
    train[overlooked_features + ['smoking']],  
    hue='smoking',  
    vars=overlooked_features,  
    diag_kind='kde',  
    plot_kws={'alpha': 0.6, 's': 40},  
    diag_kws={'fill': True}  
)  
plt.suptitle('Overlooked features', y=1.02)  
plt.show()  
  
pairs_to_plot = [  
    ('age', 'hemoglobin'),  
    ('waist(cm)', 'weight(kg)'),  
    ('triglyceride', 'systolic')  
]  
  
for x_col, y_col in pairs_to_plot:  
    g = sns.jointplot(  
        data=train,  
        x=x_col,  
        y=y_col,  
        kind='reg',  
        height=6,  
        marginal_kws={'bins': 30, 'fill': True},  
        line_kws={'color': 'darkorange'}  
    )  
    g.fig.suptitle(f'{x_col} vs {y_col}', y=1.02)  
    plt.show()  
  
pd.plotting.scatter_matrix(  
    train[overlooked_features],  
    figsize=(10, 10),  
    diagonal='kde',  
    alpha=0.5,  
    s=20  
)
```


```
plt.tight_layout()
plt.show()
```

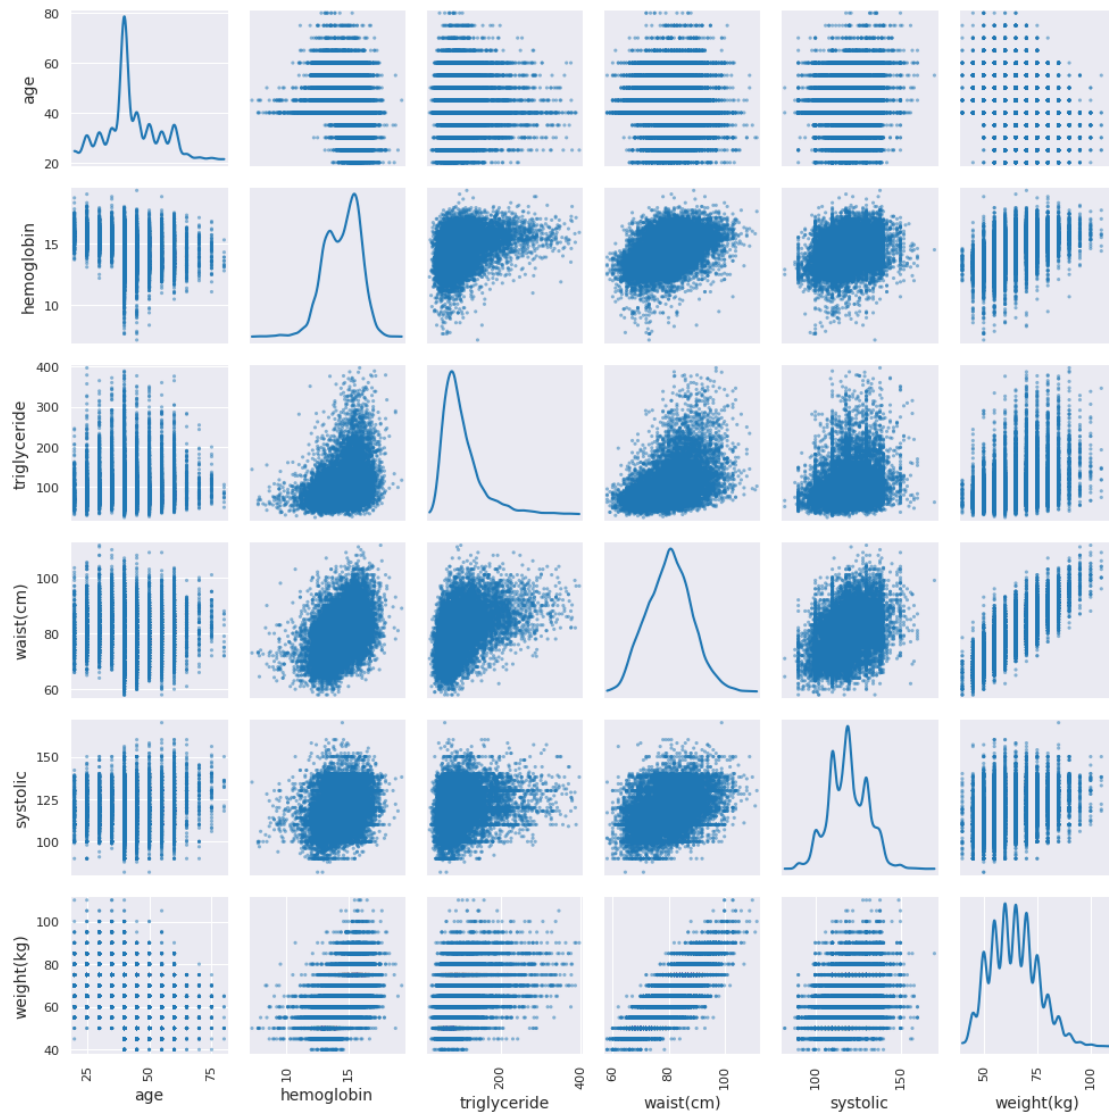
Overlooked features











## 0.0.5

EDA , , - train test  
 , , ( , , , , )  
 train test,

## 0.0.6

```
[16]: from normalization import impute_median, normalize_zscore
from feature_pruning import auto_prune_correlated_features
from cross_validation import train_test_split_stratified
from sklearn.preprocessing import OneHotEncoder
```



```

TARGET = 'smoking'
ID_COL = 'id'
overlooked_features = [column for column in train.columns if column not in
    ↳ {ID_COL, TARGET}]

X = train[overlooked_features].copy()
X_test = test[overlooked_features].copy()
y = train[TARGET].values

# Imputation
X = impute_median(X)
X_test = impute_median(X_test)

# One-Hot Encoding
categorical_cols = [column for column in overlooked_features if X[column].
    ↳ nunique() <= 10]
encoder = OneHotEncoder(sparse_output=False, drop='first')
encoder.fit(pd.concat([X[categorical_cols], X_test[categorical_cols]],
    ↳ axis=0))

X_ohe = pd.DataFrame(encoder.transform(X[categorical_cols]),
    columns=encoder.
    ↳ get_feature_names_out(categorical_cols),
    index=X.index)
X_test_ohe = pd.DataFrame(encoder.transform(X_test[categorical_cols]),
    columns=encoder.
    ↳ get_feature_names_out(categorical_cols),
    index=X_test.index)

X = pd.concat([X.drop(columns=categorical_cols), X_ohe], axis=1)
X_test = pd.concat([X_test.drop(columns=categorical_cols), X_test_ohe], axis=1)

# Drop too correlated data
df_prune = X.copy()
df_prune[TARGET] = y
corr = df_prune.drop(columns=[TARGET]).corr().abs()
upper = corr.where(np.triu(np.ones(corr.shape), k=1).astype(bool))
feature_pairs = [(c1, c2) for c1 in upper.columns for c2 in upper.index if
    ↳ upper.loc[c2, c1] > 0.8]
all_feats = X.columns.tolist()

selected = auto_prune_correlated_features(
    df=df_prune,
    target_column=TARGET,
    feature_pairs=feature_pairs,

```

```

    all_features=all_feats,
    threshold=0.8,
    verbose=True
)

X      = X[selected]
X_test = X_test[selected]

# Normalization
numeric_cols = [c for c in X.columns if c not in encoder.
                 ↪get_feature_names_out()]
X_norm, (mu, sigma) = normalize_zscore(X[numeric_cols])
X[numeric_cols]     = X_norm
X_test[numeric_cols] = (X_test[numeric_cols] - mu) / sigma

# Split
X_tr, X_val, y_tr, y_val = train_test_split_stratified(
    X, y, test_size=0.2, random_state=42
)

print(X_tr.shape, X_val.shape, X_test.shape)

```

waist(cm) vs weight(kg) → MSE\_A=0.1471, MSE\_B=0.1470 → drop f1  
 LDL vs Cholesterol → MSE\_A=0.1470, MSE\_B=0.1470 → drop f2  
 (11983, 30) (2996, 30) (10000, 30)

## 1 EDA

```

[17]: from ensemble_manual import BaggingClassifierManual
      from sklearn.ensemble import RandomForestClassifier, BaggingClassifier
      from sklearn.metrics import roc_auc_score

      bag_manual = BaggingClassifierManual(
          base_estimator=RandomForestClassifier(max_depth=3, random_state=42),
          n_estimators=50,
          max_samples=0.8,
          random_state=42
      ).fit(X_tr.values, y_tr)

      bag_sklearn = BaggingClassifier(
          estimator=RandomForestClassifier(max_depth=3, random_state=42),
          n_estimators=50,
          max_samples=0.8,
          random_state=42,
          n_jobs=-1
      ).fit(X_tr, y_tr)

```

```

p_manual = bag_manual.predict_proba(X_val.values)[: ,1]
p_sklearn = bag_sklearn.predict_proba(X_val)[: ,1]

print("Manual Bagging ROC-AUC:", roc_auc_score(y_val, p_manual))
print("Sklearn Bagging ROC-AUC:", roc_auc_score(y_val, p_sklearn))

```

Manual Bagging ROC-AUC: 0.8500621879804453  
Sklearn Bagging ROC-AUC: 0.8506538375166265

```

[18]: from ensemble_manual import GradientBoostingClassifierManual
      from sklearn.ensemble import GradientBoostingClassifier
      from sklearn.metrics import roc_auc_score
      from sklearn.tree import DecisionTreeRegressor

base_tree = DecisionTreeRegressor(max_depth=7, random_state=42)
gb_manual = GradientBoostingClassifierManual(
    base_estimator=base_tree,
    n_estimators=100, learning_rate=0.1
).fit(X_tr.values, y_tr)
p_manual = gb_manual.predict_proba(X_val.values)[: ,1]

gb_sklearn = GradientBoostingClassifier(
    n_estimators=100, learning_rate=0.1, max_depth=7, random_state=42
).fit(X_tr, y_tr)
p_skl = gb_sklearn.predict_proba(X_val)[: ,1]

print("Manual GB AUC:", roc_auc_score(y_val, p_manual))
print("Sklearn GB AUC:", roc_auc_score(y_val, p_skl))

```

Manual GB AUC: 0.8748865165325977  
Sklearn GB AUC: 0.8810832416828374

```

[19]: from sklearn.metrics import roc_auc_score
      from sklearn.tree import DecisionTreeRegressor
      from ensemble_manual import GradientBoostingClassifierManual

best_auc = -1
best_depth = None
best_model = None

for depth in range(1, 11):
    base_tree = DecisionTreeRegressor(max_depth=depth, random_state=42)
    gb_manual = GradientBoostingClassifierManual(
        base_estimator=base_tree,
        n_estimators=100,
        learning_rate=0.1
    ).fit(X_tr.values, y_tr)

```

```

p_manual = gb_manual.predict_proba(X_val.values)[: , 1]
auc = roc_auc_score(y_val, p_manual)

print(f"Depth {depth}: AUC = {auc:.4f}")

if auc > best_auc:
    best_auc = auc
    best_depth = depth
    best_model = gb_manual

print(f"\nBest manual GB depth: {best_depth} with AUC = {best_auc:.4f}")

```

```

Depth 1: AUC = 0.8384
Depth 2: AUC = 0.8511
Depth 3: AUC = 0.8603
Depth 4: AUC = 0.8650
Depth 5: AUC = 0.8695
Depth 6: AUC = 0.8732
Depth 7: AUC = 0.8749
Depth 8: AUC = 0.8742
Depth 9: AUC = 0.8729
Depth 10: AUC = 0.8713

```

Best manual GB depth: 7 with AUC = 0.8749

```

[20]: import numpy as np
from metrics import (
    accuracy_manual, precision_manual,
    recall_manual, f1_score_manual,
    roc_auc_score_manual, average_precision_score_manual
)
from sklearn.metrics import (
    accuracy_score, precision_score,
    recall_score, f1_score,
    roc_auc_score, average_precision_score
)

y_true = np.array([0, 0, 1, 1, 0, 1, 1, 0])
y_pred = np.array([0, 1, 1, 1, 0, 0, 1, 0])

y_score = y_pred.astype(float)

print("          Metric | Manual | Sklearn")
print("-----")
print(f"Accuracy          | {accuracy_manual(y_true, y_pred):7.4f} | {accuracy_score(y_true, y_pred):7.4f}")

```

```

print(f"Precision      | {precision_manual(y_true, y_pred):7.4f} |  

↳{precision_score(y_true, y_pred):7.4f}")
print(f"Recall        | {recall_manual(y_true, y_pred):7.4f} |  

↳{recall_score(y_true, y_pred):7.4f}")
print(f"F1-Score       | {f1_score_manual(y_true, y_pred):7.4f} |  

↳{f1_score(y_true, y_pred):7.4f}")
print(f"ROC-AUC        | {roc_auc_score_manual(y_true, y_score):7.4f} |  

↳{roc_auc_score(y_true, y_score):7.4f}")
print(f"PR-AUC (AP)    | {average_precision_score_manual(y_true, y_score):7.4f} |  

↳{average_precision_score(y_true, y_score):7.4f}")

```

Metric	Manual	Sklearn
Accuracy	0.7500	0.7500
Precision	0.7500	0.7500
Recall	0.7500	0.7500
F1-Score	0.7500	0.7500
ROC-AUC	0.7500	0.7500
PR-AUC (AP)	0.5193	0.6875

```

[21]: from metrics import (
        accuracy_manual, precision_manual,
        recall_manual, f1_score_manual,
        roc_auc_score_manual, average_precision_score_manual
    )
    from ensemble_manual import GradientBoostingClassifierManual
    from lightgbm import LGBMClassifier
    from xgboost import XGBClassifier
    from catboost import CatBoostClassifier
    from sklearn.ensemble import HistGradientBoostingClassifier,
    ↳GradientBoostingClassifier

    models = {
        'LightGBM': LGBMClassifier(n_estimators=100, learning_rate=0.1,
        ↳random_state=42),
        'XGBoost': XGBClassifier(n_estimators=100, learning_rate=0.1,
        ↳random_state=42, eval_metric='logloss'),
        'CatBoost': CatBoostClassifier(iterations=100, learning_rate=0.1,
        ↳random_seed=42, verbose=0),
        'SklearnHGB': HistGradientBoostingClassifier(max_iter=100, learning_rate=0.
        ↳1, random_state=42),
        'SklearnGB': GradientBoostingClassifier(n_estimators=100, learning_rate=0.
        ↳1, max_depth=3, random_state=42),
        'GB_Manual': GradientBoostingClassifierManual(
            base_estimator=DecisionTreeRegressor(max_depth=3,
        ↳random_state=42),
            n_estimators=100, learning_rate=0.1
    }

```

```

    )
}

for name, mdl in models.items():
    if name == 'GB_Manual':
        mdl.fit(X_tr.values, y_tr)
    else:
        mdl.fit(X_tr, y_tr)

metrics = [
    ('Accuracy', accuracy_manual, lambda y, y_s, y_p: y_p),
    ('Precision', precision_manual, lambda y, y_s, y_p: y_p),
    ('Recall', recall_manual, lambda y, y_s, y_p: y_p),
    ('F1', f1_score_manual, lambda y, y_s, y_p: y_p),
    ('ROC-AUC', roc_auc_score_manual, lambda y, y_s, y_p: y_s),
    ('PR-AUC', average_precision_score_manual, lambda y, y_s, y_p: y_s),
]

results = {name: {} for name in models}

for name, mdl in models.items():
    if name == 'GB_Manual':
        y_score = mdl.predict_proba(X_val.values)[: ,1]
        y_pred = mdl.predict(X_val.values)
    else:
        y_score = mdl.predict_proba(X_val)[: ,1]
        y_pred = mdl.predict(X_val)
    for mname, mfunc, selector in metrics:
        arr = selector(y_val, y_score, y_pred)
        if mname in ('ROC-AUC', 'PR-AUC'):
            val = mfunc(y_val, y_score)
        else:
            val = mfunc(y_val, y_pred)
        results[name][mname] = val

print(f"{'Model':<12}", " ".join(f"{m[0]:>9}" for m in metrics))
print("-" * (12 + 11 * len(metrics)))
for name, res in results.items():
    print(f"{name:<12}", " ".join(f"{res[m[0]]:9.4f}" for m in metrics))

best_by_metric = {}
for mname, _, _ in metrics:
    best_model = max(results, key=lambda nm: results[nm][mname])
    best_by_metric[mname] = (best_model, results[best_model][mname])

print("\nBest by metric:")
for mname, (mdl, val) in best_by_metric.items():

```

```
print(f"{mname:<9} -> {mdl} ({val:.4f})")
```

[LightGBM] [Warning] Found whitespace in feature\_names, replace with underlines  
 [LightGBM] [Info] Number of positive: 4393, number of negative: 7590  
 [LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000359 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 1099

[LightGBM] [Info] Number of data points in the train set: 11983, number of used features: 29

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.366603 -> initscore=-0.546819

[LightGBM] [Info] Start training from score -0.546819

Model	Accuracy	Precision	Recall	F1	ROC-AUC	PR-AUC
LightGBM	0.7981	0.7046	0.7732	0.7373	0.8804	0.7679
XGBoost	0.8027	0.7076	0.7869	0.7451	0.8804	0.7682
CatBoost	0.7951	0.6921	0.7942	0.7396	0.8797	0.7694
SklearnHGB	0.7977	0.7003	0.7832	0.7395	0.8785	0.7656
SklearnGB	0.7944	0.6944	0.7842	0.7365	0.8766	0.7660
GB_Manual	0.7727	0.6745	0.7341	0.7030	0.8603	0.7397

Best by metric:

Accuracy -> XGBoost (0.8027)

Precision -> XGBoost (0.7076)

Recall -> CatBoost (0.7942)

F1 -> XGBoost (0.7451)

ROC-AUC -> LightGBM (0.8804)

PR-AUC -> CatBoost (0.7694)

## OPTUNA

```
[34]: import numpy as np
import lightgbm as lgb
import optuna
from optuna.integration import LightGBMTunerCV
from sklearn.model_selection import StratifiedKFold
from lightgbm import LGBMClassifier
from sklearn.metrics import roc_auc_score

train_data = lgb.Dataset(X_tr, label=y_tr)

tuner = LightGBMTunerCV(
    params={
        'objective': 'binary',
        'metric': 'auc',
        'verbosity': -1,
        'seed': 42
    },
```

```

train_set=train_data,
folds=StratifiedKFold(n_splits=3, shuffle=True, random_state=42),
num_boost_round=500,
study=optuna.create_study(
    direction='maximize',
    pruner=optuna.pruners.SuccessiveHalvingPruner()
)
)

tuner.run()

best_params_cv = tuner.best_params
print("TunerCV best params:", best_params_cv)

final_cv = LGBMClassifier(**best_params_cv)
final_cv.fit(
    np.vstack([X_tr, X_val]),
    np.concatenate([y_tr, y_val])
)
val_preds_cv1 = final_cv.predict_proba(X_val)[: , 1]
preds_cv = final_cv.predict_proba(X_test)[: , 1]
print("TunerCV ROC-AUC on val:", roc_auc_score(y_val, val_preds_cv1))

```

[I 2025-05-17 20:38:15,141] A new study created in memory with name: no-name-8a97dee8-0552-4e3c-b0b0-01fe2b897035

0%| | 0/7 [00:00<?, ?it/s]

feature\_fraction, val\_score: -inf: 0%| | 0/7 [00:00<?,  
?it/s]/home/sergey/PycharmProjects/ML\_UNI/.venv/lib/python3.13/site-  
packages/sklearn/model\_selection/\_split.py:877: UserWarning: The groups  
parameter is ignored by StratifiedKFold  
warnings.warn(

feature\_fraction, val\_score: 0.874942: 0%| | 0/7 [00:17<?,  
?it/s]

feature\_fraction, val\_score: 0.874942: 14%|#4 | 1/7 [00:17<01:45,  
17.64s/it][I 2025-05-17 20:38:32,784] Trial 0 finished with value:  
0.8749422751595265 and parameters: {'feature\_fraction': 1.0}. Best is trial 0  
with value: 0.8749422751595265.

feature\_fraction, val\_score: 0.874942: 14%|#4 | 1/7 [00:17<01:45,  
17.64s/it]/home/sergey/PycharmProjects/ML\_UNI/.venv/lib/python3.13/site-  
packages/sklearn/model\_selection/\_split.py:877: UserWarning: The groups



```
parameter is ignored by StratifiedKFold
warnings.warn(
```

```
feature_fraction, val_score: 0.875863: 14%|#4          | 1/7 [00:19<01:45,
17.64s/it]
```

```
feature_fraction, val_score: 0.875863: 29%|##8          | 2/7 [00:19<00:41,
8.22s/it][I 2025-05-17 20:38:34,410] Trial 1 finished with value:
0.8758628486608327 and parameters: {'feature_fraction': 0.6}. Best is trial 1
with value: 0.8758628486608327.
```

```
feature_fraction, val_score: 0.875863: 29%|##8          | 2/7 [00:19<00:41,
8.22s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
warnings.warn(
```

```
feature_fraction, val_score: 0.878770: 29%|##8          | 2/7 [00:21<00:41,
8.22s/it]
```

```
feature_fraction, val_score: 0.878770: 43%|####2        | 3/7 [00:21<00:21,
5.34s/it][I 2025-05-17 20:38:36,332] Trial 2 finished with value:
0.8787701914726952 and parameters: {'feature_fraction': 0.5}. Best is trial 2
with value: 0.8787701914726952.
```

```
feature_fraction, val_score: 0.878770: 43%|####2        | 3/7 [00:21<00:21,
5.34s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
warnings.warn(
```

```
feature_fraction, val_score: 0.878770: 43%|####2        | 3/7 [00:22<00:21,
5.34s/it]
```

```
feature_fraction, val_score: 0.878770: 57%|####7        | 4/7 [00:22<00:11,
3.88s/it][I 2025-05-17 20:38:37,977] Trial 3 finished with value:
0.8784615114654778 and parameters: {'feature_fraction': 0.7}. Best is trial 2
with value: 0.8787701914726952.
```

```
feature_fraction, val_score: 0.878770: 57%|####7        | 4/7 [00:22<00:11,
3.88s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
```

```
parameter is ignored by StratifiedKFold
warnings.warn(
```

```
feature_fraction, val_score: 0.878922: 57%|#####7 | 4/7 [00:24<00:11,
3.88s/it]
```

```
feature_fraction, val_score: 0.878922: 71%|#####1 | 5/7 [00:24<00:06,
3.10s/it][I 2025-05-17 20:38:39,691] Trial 4 finished with value:
0.878922398447289 and parameters: {'feature_fraction': 0.4}. Best is trial 4
with value: 0.878922398447289.
```

```
feature_fraction, val_score: 0.878922: 71%|#####1 | 5/7 [00:24<00:06,
3.10s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
warnings.warn(
```

```
feature_fraction, val_score: 0.878922: 71%|#####1 | 5/7 [00:26<00:06,
3.10s/it]
```

```
feature_fraction, val_score: 0.878922: 86%|#####5 | 6/7 [00:26<00:02,
2.66s/it][I 2025-05-17 20:38:41,498] Trial 5 finished with value:
0.8761271854010415 and parameters: {'feature_fraction': 0.8999999999999999}.
Best is trial 4 with value: 0.878922398447289.
```

```
feature_fraction, val_score: 0.878922: 86%|#####5 | 6/7 [00:26<00:02,
2.66s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
warnings.warn(
```

```
feature_fraction, val_score: 0.878922: 86%|#####5 | 6/7 [00:28<00:02,
2.66s/it]
```

```
feature_fraction, val_score: 0.878922: 100%|#####| 7/7 [00:28<00:00,
2.42s/it][I 2025-05-17 20:38:43,421] Trial 6 finished with value:
0.8757740381613591 and parameters: {'feature_fraction': 0.8}. Best is trial 4
with value: 0.878922398447289.
```

```
feature_fraction, val_score: 0.878922: 100%|#####| 7/7 [00:28<00:00,
4.04s/it]
```

```
0%| | 0/20 [00:00<?, ?it/s]
```

```
num_leaves, val_score: 0.878922: 0%|          | 0/20 [00:00<?,
?it/s]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
num_leaves, val_score: 0.878922: 0%|          | 0/20 [00:06<?, ?it/s]
```

```
num_leaves, val_score: 0.878922: 5%|5        | 1/20 [00:06<02:07,
6.74s/it][I 2025-05-17 20:38:50,160] Trial 7 finished with value:
0.8764987784303839 and parameters: {'num_leaves': 201}. Best is trial 7 with
value: 0.8764987784303839.
```

```
num_leaves, val_score: 0.878922: 5%|5        | 1/20 [00:06<02:07,
6.74s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
num_leaves, val_score: 0.878922: 5%|5        | 1/20 [00:12<02:07,
6.74s/it]
```

```
num_leaves, val_score: 0.878922: 10%|#       | 2/20 [00:12<01:49,
6.06s/it][I 2025-05-17 20:38:55,752] Trial 8 finished with value:
0.8762219848818167 and parameters: {'num_leaves': 190}. Best is trial 7 with
value: 0.8764987784303839.
```

```
num_leaves, val_score: 0.878922: 10%|#       | 2/20 [00:12<01:49,
6.06s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
num_leaves, val_score: 0.878922: 10%|#       | 2/20 [00:17<01:49,
6.06s/it]
```

```
num_leaves, val_score: 0.878922: 15%|#5      | 3/20 [00:17<01:37,
5.74s/it][I 2025-05-17 20:39:01,114] Trial 9 finished with value:
0.8769832058872608 and parameters: {'num_leaves': 134}. Best is trial 9 with
value: 0.8769832058872608.
```

```

num_leaves, val_score: 0.878922: 15%|#5          | 3/20 [00:17<01:37,
5.74s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.878922: 15%|#5          | 3/20 [00:22<01:37,
5.74s/it]

num_leaves, val_score: 0.878922: 20%|##          | 4/20 [00:22<01:24,
5.30s/it][I 2025-05-17 20:39:05,742] Trial 10 finished with value:
0.8776814012082172 and parameters: {'num_leaves': 36}. Best is trial 10 with
value: 0.8776814012082172.

num_leaves, val_score: 0.878922: 20%|##          | 4/20 [00:22<01:24,
5.30s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.882408: 20%|##          | 4/20 [00:23<01:24,
5.30s/it]

num_leaves, val_score: 0.882408: 25%|##5         | 5/20 [00:23<00:59,
3.99s/it][I 2025-05-17 20:39:07,412] Trial 11 finished with value:
0.8824083612697905 and parameters: {'num_leaves': 7}. Best is trial 11 with
value: 0.8824083612697905.

num_leaves, val_score: 0.882408: 25%|##5         | 5/20 [00:23<00:59,
3.99s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.882408: 25%|##5         | 5/20 [00:25<00:59,
3.99s/it]

num_leaves, val_score: 0.882408: 30%|###         | 6/20 [00:25<00:43,
3.14s/it][I 2025-05-17 20:39:08,890] Trial 12 finished with value:
0.8786403281352123 and parameters: {'num_leaves': 2}. Best is trial 11 with
value: 0.8824083612697905.

```

```

num_leaves, val_score: 0.882408: 30%|###          | 6/20 [00:25<00:43,
3.14s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.882408: 30%|###          | 6/20 [00:29<00:43,
3.14s/it]

num_leaves, val_score: 0.882408: 35%|###5         | 7/20 [00:29<00:42,
3.30s/it][I 2025-05-17 20:39:12,536] Trial 13 finished with value:
0.8767084543900862 and parameters: {'num_leaves': 78}. Best is trial 11 with
value: 0.8824083612697905.

num_leaves, val_score: 0.882408: 35%|###5         | 7/20 [00:29<00:42,
3.30s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.882408: 35%|###5         | 7/20 [00:36<00:42,
3.30s/it]

num_leaves, val_score: 0.882408: 40%|####         | 8/20 [00:36<00:56,
4.74s/it][I 2025-05-17 20:39:20,366] Trial 14 finished with value:
0.8773508704770823 and parameters: {'num_leaves': 255}. Best is trial 11 with
value: 0.8824083612697905.

num_leaves, val_score: 0.882408: 40%|####         | 8/20 [00:36<00:56,
4.74s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.882408: 40%|####         | 8/20 [00:40<00:56,
4.74s/it]

num_leaves, val_score: 0.882408: 45%|####5        | 9/20 [00:40<00:49,
4.50s/it][I 2025-05-17 20:39:24,328] Trial 15 finished with value:
0.8768188039183699 and parameters: {'num_leaves': 80}. Best is trial 11 with
value: 0.8824083612697905.

```

```

num_leaves, val_score: 0.882408: 45%|####5      | 9/20 [00:40<00:49,
4.50s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.882408: 45%|####5      | 9/20 [00:46<00:49,
4.50s/it]

num_leaves, val_score: 0.882408: 50%|####5      | 10/20 [00:46<00:48,
4.84s/it][I 2025-05-17 20:39:29,923] Trial 16 finished with value:
0.8763628979767241 and parameters: {'num_leaves': 115}. Best is trial 11 with
value: 0.8824083612697905.

num_leaves, val_score: 0.882408: 50%|####5      | 10/20 [00:46<00:48,
4.84s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.884235: 50%|####5      | 10/20 [00:48<00:48,
4.84s/it]

num_leaves, val_score: 0.884235: 55%|####5      | 11/20 [00:48<00:35,
3.98s/it][I 2025-05-17 20:39:31,954] Trial 17 finished with value:
0.8842350296845054 and parameters: {'num_leaves': 5}. Best is trial 17 with
value: 0.8842350296845054.

num_leaves, val_score: 0.884235: 55%|####5      | 11/20 [00:48<00:35,
3.98s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.884235: 55%|####5      | 11/20 [00:50<00:35,
3.98s/it]

num_leaves, val_score: 0.884235: 60%|####5      | 12/20 [00:50<00:27,
3.44s/it][I 2025-05-17 20:39:34,155] Trial 18 finished with value:
0.8840966905007152 and parameters: {'num_leaves': 4}. Best is trial 17 with
value: 0.8842350296845054.

```

```

num_leaves, val_score: 0.884235: 60%|##### | 12/20 [00:50<00:27,
3.44s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.884235: 60%|##### | 12/20 [00:53<00:27,
3.44s/it]

num_leaves, val_score: 0.884235: 65%|#####5 | 13/20 [00:53<00:23,
3.39s/it][I 2025-05-17 20:39:37,423] Trial 19 finished with value:
0.8786148512681126 and parameters: {'num_leaves': 51}. Best is trial 17 with
value: 0.8842350296845054.

num_leaves, val_score: 0.884235: 65%|#####5 | 13/20 [00:54<00:23,
3.39s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.884235: 65%|#####5 | 13/20 [00:57<00:23,
3.39s/it]

num_leaves, val_score: 0.884235: 70%|##### | 14/20 [00:57<00:21,
3.52s/it][I 2025-05-17 20:39:41,263] Trial 20 finished with value:
0.8773041605633031 and parameters: {'num_leaves': 40}. Best is trial 17 with
value: 0.8842350296845054.

num_leaves, val_score: 0.884235: 70%|##### | 14/20 [00:57<00:21,
3.52s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.884235: 70%|##### | 14/20 [00:59<00:21,
3.52s/it]

num_leaves, val_score: 0.884235: 75%|#####5 | 15/20 [00:59<00:15,
3.05s/it][I 2025-05-17 20:39:43,231] Trial 21 finished with value:
0.8786403281352123 and parameters: {'num_leaves': 2}. Best is trial 17 with
value: 0.8842350296845054.

```

```

num_leaves, val_score: 0.884235: 75%|#####5 | 15/20 [00:59<00:15,
3.05s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.884235: 75%|#####5 | 15/20 [01:02<00:15,
3.05s/it]

num_leaves, val_score: 0.884235: 80%|##### | 16/20 [01:02<00:11,
2.82s/it][I 2025-05-17 20:39:45,517] Trial 22 finished with value:
0.8786403281352123 and parameters: {'num_leaves': 2}. Best is trial 17 with
value: 0.8842350296845054.

num_leaves, val_score: 0.884235: 80%|##### | 16/20 [01:02<00:11,
2.82s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.884235: 80%|##### | 16/20 [01:05<00:11,
2.82s/it]

num_leaves, val_score: 0.884235: 85%|#####5 | 17/20 [01:05<00:09,
3.05s/it][I 2025-05-17 20:39:49,084] Trial 23 finished with value:
0.8773041605633031 and parameters: {'num_leaves': 40}. Best is trial 17 with
value: 0.8842350296845054.

num_leaves, val_score: 0.884235: 85%|#####5 | 17/20 [01:05<00:09,
3.05s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.884235: 85%|#####5 | 17/20 [01:08<00:09,
3.05s/it]

num_leaves, val_score: 0.884235: 90%|##### | 18/20 [01:08<00:05,
2.99s/it][I 2025-05-17 20:39:51,934] Trial 24 finished with value:
0.879339895860662 and parameters: {'num_leaves': 24}. Best is trial 17 with
value: 0.8842350296845054.

```



```

num_leaves, val_score: 0.884235: 90%|##### | 18/20 [01:08<00:05,
2.99s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.884235: 90%|##### | 18/20 [01:13<00:05,
2.99s/it]

num_leaves, val_score: 0.884235: 95%|#####5| 19/20 [01:13<00:03,
3.44s/it][I 2025-05-17 20:39:56,439] Trial 25 finished with value:
0.8779639165796267 and parameters: {'num_leaves': 77}. Best is trial 17 with
value: 0.8842350296845054.

num_leaves, val_score: 0.884235: 95%|#####5| 19/20 [01:13<00:03,
3.44s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

num_leaves, val_score: 0.884235: 95%|#####5| 19/20 [01:16<00:03,
3.44s/it]

num_leaves, val_score: 0.884235: 100%|#####| 20/20 [01:16<00:00,
3.49s/it][I 2025-05-17 20:40:00,040] Trial 26 finished with value:
0.8778498792940117 and parameters: {'num_leaves': 62}. Best is trial 17 with
value: 0.8842350296845054.
num_leaves, val_score: 0.884235: 100%|#####| 20/20 [01:16<00:00, 3.83s/it]

0%|          | 0/10 [00:00<?, ?it/s]

bagging, val_score: 0.884235: 0%|          | 0/10 [00:00<?,
?it/s]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(

bagging, val_score: 0.884503: 0%|          | 0/10 [00:02<?, ?it/s]

bagging, val_score: 0.884503: 10%|#          | 1/10 [00:02<00:26,
2.94s/it][I 2025-05-17 20:40:02,990] Trial 27 finished with value:
0.8845030527691584 and parameters: {'bagging_fraction': 0.6991179873200065,
'bagging_freq': 1}. Best is trial 27 with value: 0.8845030527691584.

```

```
bagging, val_score: 0.884503: 10%|#          | 1/10 [00:02<00:26,
2.94s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
bagging, val_score: 0.884503: 10%|#          | 1/10 [00:05<00:26,
2.94s/it]
```

```
bagging, val_score: 0.884503: 20%|##         | 2/10 [00:05<00:22,
2.77s/it][I 2025-05-17 20:40:05,636] Trial 28 finished with value:
0.8844844905699145 and parameters: {'bagging_fraction': 0.6898969216948201,
'bagging_freq': 1}. Best is trial 27 with value: 0.8845030527691584.
```

```
bagging, val_score: 0.884503: 20%|##         | 2/10 [00:05<00:22,
2.77s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
bagging, val_score: 0.884530: 20%|##         | 2/10 [00:07<00:22,
2.77s/it]
```

```
bagging, val_score: 0.884530: 30%|###        | 3/10 [00:07<00:18,
2.58s/it][I 2025-05-17 20:40:07,983] Trial 29 finished with value:
0.8845300668570649 and parameters: {'bagging_fraction': 0.6947900450527186,
'bagging_freq': 1}. Best is trial 29 with value: 0.8845300668570649.
```

```
bagging, val_score: 0.884530: 30%|###        | 3/10 [00:07<00:18,
2.58s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
bagging, val_score: 0.885122: 30%|###        | 3/10 [00:10<00:18,
2.58s/it]
```

```
bagging, val_score: 0.885122: 40%|####       | 4/10 [00:10<00:14,
2.46s/it][I 2025-05-17 20:40:10,267] Trial 30 finished with value:
0.8851218139372605 and parameters: {'bagging_fraction': 0.6879109209649473,
'bagging_freq': 1}. Best is trial 30 with value: 0.8851218139372605.
```

```
bagging, val_score: 0.885122: 40%|####          | 4/10 [00:10<00:14,
2.46s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
bagging, val_score: 0.885122: 40%|####          | 4/10 [00:12<00:14,
2.46s/it]
```

```
bagging, val_score: 0.885122: 50%|#####         | 5/10 [00:12<00:12,
2.55s/it][I 2025-05-17 20:40:12,980] Trial 31 finished with value:
0.8849202080796433 and parameters: {'bagging_fraction': 0.6928208089515608,
'bagging_freq': 1}. Best is trial 30 with value: 0.8851218139372605.
```

```
bagging, val_score: 0.885122: 50%|#####         | 5/10 [00:12<00:12,
2.55s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
bagging, val_score: 0.885122: 50%|#####         | 5/10 [00:15<00:12,
2.55s/it]
```

```
bagging, val_score: 0.885122: 60%|#####         | 6/10 [00:15<00:09,
2.44s/it][I 2025-05-17 20:40:15,189] Trial 32 finished with value:
0.8845877304519684 and parameters: {'bagging_fraction': 0.6862709480508,
'bagging_freq': 1}. Best is trial 30 with value: 0.8851218139372605.
```

```
bagging, val_score: 0.885122: 60%|#####         | 6/10 [00:15<00:09,
2.44s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
bagging, val_score: 0.885122: 60%|#####         | 6/10 [00:17<00:09,
2.44s/it]
```

```
bagging, val_score: 0.885122: 70%|#####         | 7/10 [00:17<00:07,
2.43s/it][I 2025-05-17 20:40:17,599] Trial 33 finished with value:
0.8844496624195782 and parameters: {'bagging_fraction': 0.6713314907814028,
'bagging_freq': 1}. Best is trial 30 with value: 0.8851218139372605.
```

```
bagging, val_score: 0.885122: 70%|##### | 7/10 [00:17<00:07,
2.43s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
bagging, val_score: 0.885122: 70%|##### | 7/10 [00:19<00:07,
2.43s/it]
```

```
bagging, val_score: 0.885122: 80%|##### | 8/10 [00:19<00:04,
2.42s/it][I 2025-05-17 20:40:19,992] Trial 34 finished with value:
0.8846420140531369 and parameters: {'bagging_fraction': 0.6999031237583604,
'bagging_freq': 1}. Best is trial 30 with value: 0.8851218139372605.
```

```
bagging, val_score: 0.885122: 80%|##### | 8/10 [00:19<00:04,
2.42s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
bagging, val_score: 0.885122: 80%|##### | 8/10 [00:22<00:04,
2.42s/it]
```

```
bagging, val_score: 0.885122: 90%|##### | 9/10 [00:22<00:02,
2.36s/it][I 2025-05-17 20:40:22,244] Trial 35 finished with value:
0.8838750751563618 and parameters: {'bagging_fraction': 0.8249215395359706,
'bagging_freq': 1}. Best is trial 30 with value: 0.8851218139372605.
```

```
bagging, val_score: 0.885122: 90%|##### | 9/10 [00:22<00:02,
2.36s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
bagging, val_score: 0.885122: 90%|##### | 9/10 [00:25<00:02,
2.36s/it]
```

```
bagging, val_score: 0.885122: 100%|#####| 10/10 [00:25<00:00,
2.58s/it][I 2025-05-17 20:40:25,312] Trial 36 finished with value:
0.8831159425204245 and parameters: {'bagging_fraction': 0.5449483624916678,
'bagging_freq': 4}. Best is trial 30 with value: 0.8851218139372605.
```

bagging, val\_score: 0.885122: 100%|#####| 10/10 [00:25<00:00, 2.53s/it]

0%| | 0/3 [00:00<?, ?it/s]

feature\_fraction\_stage2, val\_score: 0.885122: 0%| | 0/3 [00:00<?,  
?it/s]/home/sergey/PycharmProjects/ML\_UNI/.venv/lib/python3.13/site-  
packages/sklearn/model\_selection/\_split.py:877: UserWarning: The groups  
parameter is ignored by StratifiedKFold  
warnings.warn(

feature\_fraction\_stage2, val\_score: 0.885122: 0%| | 0/3 [00:02<?,  
?it/s]

feature\_fraction\_stage2, val\_score: 0.885122: 33%|###3 | 1/3 [00:02<00:05,  
2.91s/it][I 2025-05-17 20:40:28,232] Trial 37 finished with value:  
0.8849850427506366 and parameters: {'feature\_fraction': 0.44800000000000006}.  
Best is trial 37 with value: 0.8849850427506366.

feature\_fraction\_stage2, val\_score: 0.885122: 33%|###3 | 1/3 [00:02<00:05,  
2.91s/it]/home/sergey/PycharmProjects/ML\_UNI/.venv/lib/python3.13/site-  
packages/sklearn/model\_selection/\_split.py:877: UserWarning: The groups  
parameter is ignored by StratifiedKFold  
warnings.warn(

feature\_fraction\_stage2, val\_score: 0.885122: 33%|###3 | 1/3 [00:05<00:05,  
2.91s/it]

feature\_fraction\_stage2, val\_score: 0.885122: 67%|#####6 | 2/3 [00:05<00:02,  
2.59s/it][I 2025-05-17 20:40:30,601] Trial 38 finished with value:  
0.8846952706699129 and parameters: {'feature\_fraction': 0.48000000000000004}.  
Best is trial 37 with value: 0.8849850427506366.

feature\_fraction\_stage2, val\_score: 0.885122: 67%|#####6 | 2/3 [00:05<00:02,  
2.59s/it]/home/sergey/PycharmProjects/ML\_UNI/.venv/lib/python3.13/site-  
packages/sklearn/model\_selection/\_split.py:877: UserWarning: The groups  
parameter is ignored by StratifiedKFold  
warnings.warn(

feature\_fraction\_stage2, val\_score: 0.885122: 67%|#####6 | 2/3 [00:07<00:02,  
2.59s/it]

feature\_fraction\_stage2, val\_score: 0.885122: 100%|#####| 3/3 [00:07<00:00,

```
2.46s/it][I 2025-05-17 20:40:32,908] Trial 39 finished with value:
0.8851218139372605 and parameters: {'feature_fraction': 0.41600000000000004}.
Best is trial 39 with value: 0.8851218139372605.
feature_fraction_stage2, val_score: 0.885122: 100%|#####| 3/3 [00:07<00:00,
2.53s/it]
```

```
0%|          | 0/20 [00:00<?, ?it/s]
```

```
regularization_factors, val_score: 0.885122: 0%|          | 0/20 [00:00<?,
?it/s]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
regularization_factors, val_score: 0.885122: 0%|          | 0/20 [00:02<?,
?it/s]
```

```
regularization_factors, val_score: 0.885122: 5%|5          | 1/20 [00:02<00:56,
2.96s/it][I 2025-05-17 20:40:35,876] Trial 40 finished with value:
0.8845594646692428 and parameters: {'lambda_l1': 0.006715080016653908,
'lambda_l2': 7.033668784265104e-06}. Best is trial 40 with value:
0.8845594646692428.
```

```
regularization_factors, val_score: 0.885122: 5%|5          | 1/20 [00:02<00:56,
2.96s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
regularization_factors, val_score: 0.886033: 5%|5          | 1/20 [00:05<00:56,
2.96s/it]
```

```
regularization_factors, val_score: 0.886033: 10%|#         | 2/20 [00:05<00:46,
2.58s/it][I 2025-05-17 20:40:38,185] Trial 41 finished with value:
0.8860328124277815 and parameters: {'lambda_l1': 5.962957689728475e-08,
'lambda_l2': 7.364494123592756}. Best is trial 41 with value:
0.8860328124277815.
```

```
regularization_factors, val_score: 0.886033: 10%|#         | 2/20 [00:05<00:46,
2.58s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

regularization\_factors, val\_score: 0.886707: 10%|# | 2/20 [00:07<00:46, 2.58s/it]

regularization\_factors, val\_score: 0.886707: 15%|#5 | 3/20 [00:07<00:41, 2.42s/it] [I 2025-05-17 20:40:40,420] Trial 42 finished with value: 0.8867068250859201 and parameters: {'lambda\_l1': 1.4619627113024688e-08, 'lambda\_l2': 7.830054025447685}. Best is trial 42 with value: 0.8867068250859201.

regularization\_factors, val\_score: 0.886707: 15%|#5 | 3/20 [00:07<00:41, 2.42s/it]/home/sergey/PycharmProjects/ML\_UNI/.venv/lib/python3.13/site-packages/sklearn/model\_selection/\_split.py:877: UserWarning: The groups parameter is ignored by StratifiedKFold  
warnings.warn(

regularization\_factors, val\_score: 0.886707: 15%|#5 | 3/20 [00:09<00:41, 2.42s/it]

regularization\_factors, val\_score: 0.886707: 20%|## | 4/20 [00:09<00:38, 2.39s/it] [I 2025-05-17 20:40:42,770] Trial 43 finished with value: 0.8860668823493608 and parameters: {'lambda\_l1': 1.7036489088261837e-08, 'lambda\_l2': 8.885164541968038}. Best is trial 42 with value: 0.8867068250859201.

regularization\_factors, val\_score: 0.886707: 20%|## | 4/20 [00:09<00:38, 2.39s/it]/home/sergey/PycharmProjects/ML\_UNI/.venv/lib/python3.13/site-packages/sklearn/model\_selection/\_split.py:877: UserWarning: The groups parameter is ignored by StratifiedKFold  
warnings.warn(

regularization\_factors, val\_score: 0.886707: 20%|## | 4/20 [00:13<00:38, 2.39s/it]

regularization\_factors, val\_score: 0.886707: 25%|##5 | 5/20 [00:13<00:40, 2.70s/it] [I 2025-05-17 20:40:46,012] Trial 44 finished with value: 0.886283103169164 and parameters: {'lambda\_l1': 1.0078748015971085e-08, 'lambda\_l2': 8.327798351391525}. Best is trial 42 with value: 0.8867068250859201.

regularization\_factors, val\_score: 0.886707: 25%|##5 | 5/20 [00:13<00:40, 2.70s/it]/home/sergey/PycharmProjects/ML\_UNI/.venv/lib/python3.13/site-

```
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 25%|##5          | 5/20 [00:15<00:40,
2.70s/it]
```

```
regularization_factors, val_score: 0.886707: 30%|###          | 6/20 [00:15<00:36,
2.58s/it][I 2025-05-17 20:40:48,363] Trial 45 finished with value:
0.8860363629508065 and parameters: {'lambda_l1': 1.1208101167456596e-08,
'lambda_l2': 5.271405600277059}. Best is trial 42 with value:
0.8867068250859201.
```

```
regularization_factors, val_score: 0.886707: 30%|###          | 6/20 [00:15<00:36,
2.58s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 30%|###          | 6/20 [00:18<00:36,
2.58s/it]
```

```
regularization_factors, val_score: 0.886707: 35%|###5         | 7/20 [00:18<00:33,
2.57s/it][I 2025-05-17 20:40:50,919] Trial 46 finished with value:
0.8859480836967585 and parameters: {'lambda_l1': 1.2715263505638996e-08,
'lambda_l2': 7.2237717887996595}. Best is trial 42 with value:
0.8867068250859201.
```

```
regularization_factors, val_score: 0.886707: 35%|###5         | 7/20 [00:18<00:33,
2.57s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 35%|###5         | 7/20 [00:20<00:33,
2.57s/it]
```

```
regularization_factors, val_score: 0.886707: 40%|####         | 8/20 [00:20<00:30,
2.52s/it][I 2025-05-17 20:40:53,314] Trial 47 finished with value:
0.8858095496574535 and parameters: {'lambda_l1': 1.141977350078899e-08,
'lambda_l2': 6.947767156047918}. Best is trial 42 with value:
0.8867068250859201.
```



```
regularization_factors, val_score: 0.886707: 40%|####          | 8/20 [00:20<00:30,
2.52s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
```

```
warnings.warn(
num_leaves, val_score: 0.884235: 50%|#####          | 10/20 [05:14<05:14, 31.40s/it]
```

```
regularization_factors, val_score: 0.886707: 40%|####          | 8/20 [00:23<00:30,
2.52s/it]
```

```
regularization_factors, val_score: 0.886707: 45%|#####5      | 9/20 [00:23<00:30,
2.80s/it][I 2025-05-17 20:40:56,744] Trial 48 finished with value:
0.8856819625419883 and parameters: {'lambda_l1': 1.4022102814516587e-08,
'lambda_l2': 9.067745633148425}. Best is trial 42 with value:
0.8867068250859201.
```

```
regularization_factors, val_score: 0.886707: 45%|#####5      | 9/20 [00:23<00:30,
2.80s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
```

```
warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 45%|#####5      | 9/20 [00:26<00:30,
2.80s/it]
```

```
regularization_factors, val_score: 0.886707: 50%|#####          | 10/20
[00:26<00:26, 2.67s/it][I 2025-05-17 20:40:59,125] Trial 49 finished with
value: 0.8857398635018691 and parameters: {'lambda_l1': 1.1367317720112154e-08,
'lambda_l2': 6.942880635441359}. Best is trial 42 with value:
0.8867068250859201.
```

```
regularization_factors, val_score: 0.886707: 50%|#####          | 10/20
[00:26<00:26,
2.67s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
```

```
warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 50%|#####          | 10/20
[00:28<00:26, 2.67s/it]
```

```
regularization_factors, val_score: 0.886707: 55%|#####5      | 11/20
```

```
[00:28<00:23, 2.59s/it][I 2025-05-17 20:41:01,531] Trial 50 finished with
value: 0.885872909766097 and parameters: {'lambda_11': 1.2766782408964144e-07,
'lambda_12': 7.001281645161147}. Best is trial 42 with value:
0.8867068250859201.
```

```
regularization_factors, val_score: 0.886707: 55%|#####5 | 11/20
[00:28<00:23,
2.59s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 55%|#####5 | 11/20
[00:31<00:23, 2.59s/it]
```

```
regularization_factors, val_score: 0.886707: 60%|##### | 12/20
[00:31<00:20, 2.57s/it][I 2025-05-17 20:41:04,061] Trial 51 finished with
value: 0.8860076453514637 and parameters: {'lambda_11': 1.2972037957493148e-08,
'lambda_12': 8.651445415945725}. Best is trial 42 with value:
0.8867068250859201.
```

```
regularization_factors, val_score: 0.886707: 60%|##### | 12/20
[00:31<00:20,
2.57s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 60%|##### | 12/20
[00:34<00:20, 2.57s/it]
```

```
regularization_factors, val_score: 0.886707: 65%|#####5 | 13/20
[00:34<00:19, 2.84s/it][I 2025-05-17 20:41:07,516] Trial 52 finished with
value: 0.8856652370341163 and parameters: {'lambda_11': 6.695290538482802e-08,
'lambda_12': 4.420357418925196}. Best is trial 42 with value:
0.8867068250859201.
```

```
regularization_factors, val_score: 0.886707: 65%|#####5 | 13/20
[00:34<00:19,
2.84s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 65%|#####5 | 13/20
[00:37<00:19, 2.84s/it]
```

```
regularization_factors, val_score: 0.886707: 70%|##### | 14/20
[00:37<00:16, 2.74s/it][I 2025-05-17 20:41:10,022] Trial 53 finished with
value: 0.8850525173084157 and parameters: {'lambda_l1': 7.275344039890602e-07,
'lambda_l2': 0.27470530459445136}. Best is trial 42 with value:
0.8867068250859201.
```

```
regularization_factors, val_score: 0.886707: 70%|##### | 14/20
[00:37<00:16,
2.74s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 70%|##### | 14/20
[00:39<00:16, 2.74s/it]
```

```
regularization_factors, val_score: 0.886707: 75%|#####5 | 15/20
[00:39<00:13, 2.61s/it][I 2025-05-17 20:41:12,318] Trial 54 finished with
value: 0.8851680992478504 and parameters: {'lambda_l1': 4.872838664016751e-06,
'lambda_l2': 0.03989342804460906}. Best is trial 42 with value:
0.8867068250859201.
```

```
regularization_factors, val_score: 0.886707: 75%|#####5 | 15/20
[00:39<00:13,
2.61s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 75%|#####5 | 15/20
[00:41<00:13, 2.61s/it]
```

```
regularization_factors, val_score: 0.886707: 80%|##### | 16/20
[00:41<00:10, 2.53s/it][I 2025-05-17 20:41:14,660] Trial 55 finished with
value: 0.885385092056699 and parameters: {'lambda_l1': 2.1670936752144563e-08,
'lambda_l2': 0.18154257885372585}. Best is trial 42 with value:
0.8867068250859201.
```

```
regularization_factors, val_score: 0.886707: 80%|##### | 16/20
[00:41<00:10,
2.53s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 80%|##### | 16/20
[00:44<00:10, 2.53s/it]
```

```
regularization_factors, val_score: 0.886707: 85%|#####5 | 17/20
[00:44<00:08, 2.69s/it][I 2025-05-17 20:41:17,736] Trial 56 finished with
value: 0.8845010416169853 and parameters: {'lambda_11': 6.89902659229063e-07,
'lambda_12': 0.28950684709744046}. Best is trial 42 with value:
0.8867068250859201.
```

```
regularization_factors, val_score: 0.886707: 85%|#####5 | 17/20
[00:44<00:08,
2.69s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 85%|#####5 | 17/20
[00:47<00:08, 2.69s/it]
```

```
regularization_factors, val_score: 0.886707: 90%|##### | 18/20
[00:47<00:05, 2.60s/it][I 2025-05-17 20:41:20,113] Trial 57 finished with
value: 0.8841830736607167 and parameters: {'lambda_11': 8.932510071239893,
'lambda_12': 0.4725059790304198}. Best is trial 42 with value:
0.8867068250859201.
```

```
regularization_factors, val_score: 0.886707: 90%|##### | 18/20
[00:47<00:05,
2.60s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
  warnings.warn(
```

```
regularization_factors, val_score: 0.886707: 90%|##### | 18/20
[00:49<00:05, 2.60s/it]
```

```
regularization_factors, val_score: 0.886707: 95%|#####5 | 19/20
```

[00:49<00:02, 2.53s/it] [I 2025-05-17 20:41:22,492] Trial 58 finished with value: 0.8847930141767093 and parameters: {'lambda\_l1': 1.9404593240889336e-07, 'lambda\_l2': 0.0009288159394981387}. Best is trial 42 with value: 0.8867068250859201.

regularization\_factors, val\_score: 0.886707: 95%|#####5| 19/20  
[00:49<00:02, 2.53s/it]/home/sergey/PycharmProjects/ML\_UNI/.venv/lib/python3.13/site-packages/sklearn/model\_selection/\_split.py:877: UserWarning: The groups parameter is ignored by StratifiedKFold  
warnings.warn(

regularization\_factors, val\_score: 0.886707: 95%|#####5| 19/20  
[00:52<00:02, 2.53s/it]

regularization\_factors, val\_score: 0.886707: 100%|#####| 20/20  
[00:52<00:00, 2.50s/it] [I 2025-05-17 20:41:24,921] Trial 59 finished with value: 0.8851218139372605 and parameters: {'lambda\_l1': 1.0065620632728834e-08, 'lambda\_l2': 1.1749010343263647e-08}. Best is trial 42 with value: 0.8867068250859201.

regularization\_factors, val\_score: 0.886707: 100%|#####| 20/20  
[00:52<00:00, 2.60s/it]

min\_child\_samples, val\_score: 0.886707: 0%| | 0/5 [00:00<?, ?it/s]/home/sergey/PycharmProjects/ML\_UNI/.venv/lib/python3.13/site-packages/sklearn/model\_selection/\_split.py:877: UserWarning: The groups parameter is ignored by StratifiedKFold  
warnings.warn(

min\_child\_samples, val\_score: 0.886707: 20%|## | 1/5 [00:02<00:10, 2.51s/it] [I 2025-05-17 20:41:27,432] Trial 60 finished with value: 0.8858991029053596 and parameters: {'min\_child\_samples': 50}. Best is trial 60 with value: 0.8858991029053596.

min\_child\_samples, val\_score: 0.886707: 20%|## | 1/5 [00:02<00:10, 2.51s/it]/home/sergey/PycharmProjects/ML\_UNI/.venv/lib/python3.13/site-packages/sklearn/model\_selection/\_split.py:877: UserWarning: The groups parameter is ignored by StratifiedKFold  
warnings.warn(

min\_child\_samples, val\_score: 0.886707: 40%|#### | 2/5 [00:05<00:08, 2.67s/it] [I 2025-05-17 20:41:30,220] Trial 61 finished with value: 0.8853823713896474 and parameters: {'min\_child\_samples': 100}. Best is trial 60 with value: 0.8858991029053596.

min\_child\_samples, val\_score: 0.886707: 40%|#### | 2/5 [00:05<00:08, 2.67s/it]/home/sergey/PycharmProjects/ML\_UNI/.venv/lib/python3.13/site-packages/sklearn/model\_selection/\_split.py:877: UserWarning: The groups parameter is ignored by StratifiedKFold  
warnings.warn(

min\_child\_samples, val\_score: 0.886707: 60%|##### | 3/5 [00:07<00:04,

```

2.27s/it][I 2025-05-17 20:41:32,009] Trial 62 finished with value:
0.886152051593907 and parameters: {'min_child_samples': 5}. Best is trial 62
with value: 0.886152051593907.
min_child_samples, val_score: 0.886707: 60%|##### | 3/5 [00:07<00:04,
2.27s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(
min_child_samples, val_score: 0.886707: 80%|##### | 4/5 [00:08<00:02,
2.08s/it][I 2025-05-17 20:41:33,796] Trial 63 finished with value:
0.8862019982469884 and parameters: {'min_child_samples': 10}. Best is trial 63
with value: 0.8862019982469884.
min_child_samples, val_score: 0.886707: 80%|##### | 4/5 [00:08<00:02,
2.08s/it]/home/sergey/PycharmProjects/ML_UNI/.venv/lib/python3.13/site-
packages/sklearn/model_selection/_split.py:877: UserWarning: The groups
parameter is ignored by StratifiedKFold
    warnings.warn(
min_child_samples, val_score: 0.886707: 100%|##### | 5/5 [00:10<00:00,
1.94s/it][I 2025-05-17 20:41:35,496] Trial 64 finished with value:
0.8859250948004022 and parameters: {'min_child_samples': 25}. Best is trial 63
with value: 0.8862019982469884.
min_child_samples, val_score: 0.886707: 100%|##### | 5/5 [00:10<00:00,
2.11s/it]

TunerCV best params: {'objective': 'binary', 'metric': 'auc', 'verbosity': -1,
'seed': 42, 'feature_pre_filter': False, 'lambda_l1': 1.4619627113024688e-08,
'lambda_l2': 7.830054025447685, 'num_leaves': 5, 'feature_fraction': 0.4,
'bagging_fraction': 0.6879109209649473, 'bagging_freq': 1, 'min_child_samples':
20}
TunerCV ROC-AUC on val: 0.8854517553709109

```

```

[ ]: from sklearn.model_selection import cross_val_score

def objective(trial):
    params = {
        'n_estimators': trial.suggest_int('n_estimators', 50, 500),
        'max_depth': trial.suggest_int('max_depth', 3, 12),
        'num_leaves': trial.suggest_int('num_leaves', 20, 200),
        'learning_rate': trial.suggest_loguniform('learning_rate', 1e-3, 1.
    ↪0),
        'min_child_samples': trial.suggest_int('min_child_samples', 5, 100),
        'subsample': trial.suggest_uniform('subsample', 0.5, 1.0),
        'colsample_bytree': trial.suggest_uniform('colsample_bytree', 0.5, 1.
    ↪0),
        'reg_alpha': trial.suggest_loguniform('reg_alpha', 1e-8, 10.0),
        'reg_lambda': trial.suggest_loguniform('reg_lambda', 1e-8, 10.
    ↪0),
        'random_state': 42,

```

```

        'objective':          'binary',
        'verbosity':          -1
    }
    model = LGBMClassifier(**params)
    cv = StratifiedKFold(n_splits=3, shuffle=True, random_state=42)
    scores = cross_val_score(model, X_tr, y_tr, cv=cv, scoring='roc_auc',
↪n_jobs=-1)
    return scores.mean()

study = optuna.create_study(direction='maximize')
study.optimize(objective, n_trials=50, timeout=600)

best_params_cv2 = study.best_params
print("Classic Optuna best params:", best_params_cv2)

final2 = LGBMClassifier(**best_params_cv2)
final2.fit(
    np.vstack([X_tr, X_val]),
    np.concatenate([y_tr, y_val])
)
val_preds_cv2 = final2.predict_proba(X_val)[: , 1]
preds_cv2 = final2.predict_proba(X_test)[: , 1]

print("TunerCV ROC-AUC on val:", roc_auc_score(y_val, val_preds_cv1))
print("Classic Optuna ROC-AUC on val:", roc_auc_score(y_val, val_preds_cv2))

```

```

[35]: import pandas as pd

sample = pd.read_csv('mai-ml-lab-2/sample_submission.csv')

sample['smoking'] = preds_cv

sample.to_csv('submission.csv', index=False)

print("submission.csv      :", sample.head())

```

```

submission.csv      :      id  smoking
0  15000  0.044976
1  15001  0.820014
2  15002  0.206059
3  15003  0.720090
4  15004  0.015359

```